

# Chapter 15

## Adaptive Thermal Intelligence

*From Static Metamaterials to the Dynamic Mode-Path Matrix*

Thermal Photonics: A Unified Framework for Engineering Heat Flux  
Book 3 of the Open-Source Engineering Optics Trilogy

Manuscript build: tpd-ch15-v4

### Learning Objectives

- LO-15.1** Explain the physical principles of optical refrigeration (anti-Stokes fluorescence) and evaluate its feasibility for below-ambient semiconductor cooling using the Mode-Path Matrix (MPM).
- LO-15.2** Classify phase-change reconfigurable metamaterials (VO<sub>2</sub>, GST) by switching speed, contrast ratio, and volatility, and assess their impact on dynamic thermal network topology.
- LO-15.3** Describe how machine learning inverse design optimizes thermal metamaterial structures and how it integrates with the MPM workflow.
- LO-15.4** Design an adaptive thermal network with sensor-driven conductance adjustment and evaluate its stability using closed-loop control theory.
- LO-15.5** Assess the feasibility and limitations of thermal computing concepts (heat-based logic gates, thermal memory) against established semiconductor technology benchmarks.
- LO-15.6** Articulate how the three books of the Open-Source Engineering Optics Trilogy form a complete photonic engineering design stack covering the full photon energy lifecycle.

## Abbreviations and Nomenclature

### Acronyms

| Abbreviation | Meaning                                                              |
|--------------|----------------------------------------------------------------------|
| ASF          | Anti-Stokes fluorescence                                             |
| COP          | Coefficient of performance                                           |
| DRL          | Deep reinforcement learning                                          |
| DR           | Design Rule                                                          |
| EMT          | Effective Medium Theory                                              |
| FEA          | Finite Element Analysis                                              |
| GAN          | Generative Adversarial Network                                       |
| GST          | Ge <sub>2</sub> Sb <sub>2</sub> Te <sub>5</sub> (phase-change alloy) |
| HBM          | High Bandwidth Memory                                                |
| LWIR         | Long-wave infrared (8–14 $\mu\text{m}$ )                             |
| MIT          | Metal–Insulator Transition                                           |

| Abbreviation | Meaning                          |
|--------------|----------------------------------|
| ML           | Machine Learning                 |
| MPM          | Mode-Path Matrix                 |
| MPC          | Model Predictive Control         |
| PID          | Proportional-Integral-Derivative |
| RTD          | Resistance Temperature Detector  |
| TEC          | Thermoelectric cooler            |
| TIM          | Thermal Interface Material       |
| TPV          | Thermophotovoltaic               |
| TRL          | Technology Readiness Level       |
| VAE          | Variational Autoencoder          |
| WE           | Worked Example                   |

## Principal Symbols

| Symbol               | Meaning                                               | SI Unit                             |
|----------------------|-------------------------------------------------------|-------------------------------------|
| $\Gamma$             | Mode dominance ratio $G_{\text{rad}}/G_{\text{cond}}$ | —                                   |
| $G_{\text{cond}}$    | Conduction-mode thermal conductance                   | W/K                                 |
| $G_{\text{rad}}$     | Radiation-mode thermal conductance                    | W/K                                 |
| $G_{\text{active}}$  | Active edge effective conductance                     | W/K                                 |
| $\kappa$             | Thermal conductivity                                  | W/(m·K)                             |
| $\varepsilon$        | Emissivity                                            | —                                   |
| $\sigma_{\text{SB}}$ | Stefan–Boltzmann constant                             | W/(m <sup>2</sup> ·K <sup>4</sup> ) |
| $\alpha$             | Seebeck coefficient                                   | V/K                                 |
| $\eta_{\text{ext}}$  | External quantum efficiency                           | —                                   |
| $P_{\text{cool}}$    | Cooling power                                         | W                                   |
| $P_{\text{elec}}$    | Electrical input power                                | W                                   |
| COP                  | Coefficient of performance                            | —                                   |
| $\tau_{\text{th}}$   | Thermal time constant                                 | s                                   |
| $C$                  | Thermal capacitance                                   | J/K                                 |
| $T_{\text{MIT}}$     | Metal–insulator transition temperature                | K (°C)                              |

## Subscript and Superscript Conventions

| Notation        | Convention                                         |
|-----------------|----------------------------------------------------|
| cond, rad       | Conduction and radiation modes                     |
| active, passive | Pumped and unpumped contributions                  |
| ins, met        | Insulating and metallic phases (VO <sub>2</sub> )  |
| amorph, cryst   | Amorphous and crystalline phases (GST)             |
| switch          | Pertaining to switchable or reconfigurable element |
| on, off         | High-conductance and low-conductance states        |
| opt             | Optimized (ML-derived) value                       |
| target          | Design objective value                             |

*Design rules for this chapter follow the DR-15F.n numbering convention (F = Future). Topological design rules inherited from Chapter 4 retain the DR-T.n prefix. Worked examples use WE-15.n. Learning objectives use LO-15.n.*

## Chapter Roadmap and Deliverables

This chapter is organized around a single engineering deliverable:

**A technology readiness assessment of active and adaptive thermal control techniques—evaluated through the MPM framework—that identifies which emerging methods are nearest to semiconductor production and which remain research-grade.**

To help you navigate, Figure 15.1 shows the chapter “spine” (workflow) and how each section feeds into the technology readiness assessment.

**Figure 15.1: Chapter 15 Roadmap**




Figure 15.1: Chapter 15 roadmap. Four emerging technology domains (§15.1–§15.4) are evaluated through the MPM framework, producing a technology readiness assessment. Thermal computing (§15.5) represents the speculative horizon, while the trilogy coda (§15.6) closes the three-book narrative arc. Color coding follows the TPD scheme: blue (conduction), red (radiation), purple (multiphysics), green (metamaterial/design), teal (future/active).

**Why this chapter is considered.** Chapters 1–14 have developed a comprehensive, systematic toolkit for thermal engineering based on the MPM framework. Every tool presented so far—transformation thermotics, effective medium theory, via arrays, spectral emitters, multiphysics cloaks, and quantum thermal diagnostics—operates in a fundamentally *passive* paradigm: the thermal conductances, emissivities, and network topologies are fixed at fabrication time and cannot adapt to changing operating conditions.

**Pain Point**

Production thermal systems face dynamic operating conditions—varying power loads, ambient drift, process transients—yet the metamaterial designs from Chapters 5–9 are frozen at fabrication. When the critical path shifts (DR-T.6), the passive design may no longer address the true bottleneck. An *active* thermal management paradigm is needed: one that can reconfigure network topology, adjust conductances in real time, and exploit emerging computational tools for design optimization.

**Mode-Path Matrix (MPM) Refresher**

Chapter 15 depends on the MPM more than any other chapter. For readers arriving here or needing a quick reminder, the core concepts are:

1. **Nodes** = thermal subsystems (dies, spreaders, heat sinks, ambient).
2. **Edges** = transport channels connecting nodes, each carrying conduction and/or radiation.
3. **Edge conductances:**  $G_{\text{cond}} = \kappa A/L$  (conduction);  $G_{\text{rad}} = 4\epsilon\sigma_{\text{SB}}T^3A$  (linearized radiation).
4. **Mode dominance ratio:**  $\Gamma = G_{\text{rad}}/G_{\text{cond}}$ .
  - $\Gamma < 0.1$ : conduction-dominated  $\rightarrow$  use Ch. 5–7 tools.
  - $\Gamma > 10$ : radiation-dominated  $\rightarrow$  use Ch. 8–10 tools.
  - $0.1 \leq \Gamma \leq 10$ : mixed-mode  $\rightarrow$  use Ch. 9 multiphysics coupling.
5. **Critical path** = the path with the smallest total conductance (largest thermal resistance); this is where the temperature drop is greatest and where design effort has the most impact (DR-T.3).
6. **Critical-path shift (DR-T.6):** Modifying one edge can migrate the bottleneck to a different path—*this is the central hazard that Chapter 15 addresses*. In adaptive systems, the network topology itself changes with time, so the critical path must be re-evaluated after every intervention.

*Full derivation:* Ch. 4 (§04.2–04.4). *Linearization validity:* DR-T.1b. *Fill-out recipe:* Ch. 4, §04.4.2.

This final chapter surveys four active and adaptive technologies that extend the MPM framework from static to dynamic thermal control: externally pumped cooling (§15.1), reconfigurable metamaterials (§15.2), machine learning-assisted design (§15.3), and adaptive thermal networks (§15.4). For each, we evaluate the current technology readiness, identify the MPM integration strategy, and assess the timeline to semiconductor production relevance. Section 15.5 examines the speculative frontier of thermal computing, and Section 15.6 closes the trilogy by connecting the three books into a unified photonic engineering design stack.

## 15.1 Externally Pumped Thermal Control

All thermal devices in Chapters 5–10 are thermodynamically passive: they redirect or concentrate heat flux but cannot remove entropy from a system. Achieving temperatures below the ambient or heat-sink temperature requires an externally pumped mechanism—a heat pump in

the thermodynamic sense. This section examines two pumped-cooling technologies through the MPM lens.

### 15.1.1 Thermoelectric Cooling: The Baseline

Thermoelectric coolers (TECs) based on the Peltier effect are the most mature actively pumped thermal devices in semiconductor manufacturing. A TEC module placed on an MPM edge acts as a negative thermal resistance—it pumps heat against the temperature gradient. The cooling power is:

$$P_{\text{cool}} = \alpha IT_{\text{cold}} - \frac{1}{2}I^2 R_{\text{TEC}} - G_{\text{TEC}}\Delta T \quad (15.1)$$

where  $\alpha$  is the Seebeck coefficient,  $I$  is the drive current,  $R_{\text{TEC}}$  is the module electrical resistance, and  $G_{\text{TEC}}$  is the parasitic thermal conductance. The coefficient of performance is:

$$\text{COP} = \frac{P_{\text{cool}}}{P_{\text{electrical}}} = \frac{P_{\text{cool}}}{\alpha I \Delta T + I^2 R_{\text{TEC}}} \quad (15.2)$$

For semiconductor applications, single-stage  $\text{Bi}_2\text{Te}_3$  TECs achieve  $\Delta T_{\text{max}} \approx 70$  K with  $\text{COP} \approx 0.3\text{--}1.0$  depending on the operating point. Multi-stage cascaded TECs can reach  $\Delta T_{\text{max}} \approx 130$  K but with significantly reduced COP ( $< 0.1$ ).

**Figure 15.2: Active Cooling Mechanisms for Mode-Path Matrix Integration**




Figure 15.2: Active cooling mechanisms for MPM integration. **(a)** Thermoelectric cooler (TEC): Peltier effect pumps heat from cold junction to hot junction via charge carriers; conduction-dominated ( $\Gamma \ll 0.1$ ). **(b)** Optical refrigeration via anti-Stokes fluorescence (ASF): pump photons at  $\lambda_{\text{pump}}$  excite  $\text{Yb}^{3+}$  ions; mean fluorescence at shorter  $\lambda_{\text{fluor}}$  carries away thermal energy; radiation-dominated ( $\Gamma \gg 10$ ). Both mechanisms introduce active edges in the thermal network with effective negative resistance. Insets show the MPM edge representation for each mechanism.

#### Conduction Mode (Diffusionics)

In the MPM, a TEC edge is conduction-dominated ( $\Gamma \ll 0.1$ ) because the Peltier heat transport is a bulk carrier phenomenon. The TEC modifies the effective conductance of its edge:  $G_{\text{eff}} = G_{\text{TEC}} + P_{\text{cool}}/\Delta T$ , which can be made negative (net heat removal against the gradient). DR-T.5 (conductance matching) still applies: the TEC must be matched to adjacent elements to avoid bottlenecks.

### 15.1.2 Optical Refrigeration: Anti-Stokes Fluorescence

Optical refrigeration exploits anti-Stokes fluorescence (ASF) to achieve vibrational cooling without mechanical parts or electrical contacts. The principle is elegant: a laser at wavelength  $\lambda_{\text{pump}}$  excites a fluorescent medium (typically  $\text{Yb}^{3+}$ -doped glass or crystal), which emits at a shorter mean wavelength  $\lambda_{\text{fluor}} < \lambda_{\text{pump}}$ . Each emitted photon carries away more energy than was absorbed, and the energy difference is extracted from the thermal phonon bath.

The net cooling power per unit volume is:

$$p_{\text{cool}} = \eta_{\text{ext}} \eta_{\text{abs}} I_{\text{pump}} \left( \frac{\lambda_{\text{pump}}}{\lambda_{\text{fluor}}} - 1 \right) \quad (15.3)$$

where  $\eta_{\text{ext}}$  is the external quantum efficiency (fraction of absorbed photons that produce escaping fluorescence),  $\eta_{\text{abs}}$  is the absorption efficiency, and  $I_{\text{pump}}$  is the pump intensity. The term in parentheses is the fractional anti-Stokes shift, typically 1–3% for  $\text{Yb}^{3+}$  systems.

Table 15.4: Key parameters for optical refrigeration media.

| Material          | $\lambda_{\text{pump}}$ (nm) | $\eta_{\text{ext}}$ | $\Delta T_{\text{max}}$ (K) | Status   |
|-------------------|------------------------------|---------------------|-----------------------------|----------|
| Yb:YLF crystal    | 1020                         | 0.996               | 174                         | Lab demo |
| Yb:ZBLAN glass    | 1030                         | 0.990               | 70                          | Lab demo |
| Yb:YAG            | 1030                         | $> 0.99$            | TBD                         | Research |
| CdSe quantum dots | 520                          | 0.80                | 30                          | Emerging |

The critical requirement is  $\eta_{\text{ext}} > \eta_{\text{ext,min}}$ , below which parasitic absorption causes net heating rather than cooling. For  $\text{Yb}^{3+}$  at  $\lambda_{\text{pump}}/\lambda_{\text{fluor}} \approx 0.98$ , the minimum external quantum efficiency is:

$$\eta_{\text{ext,min}} = \frac{\lambda_{\text{fluor}}}{\lambda_{\text{pump}}} \approx \frac{1}{1 + 0.02} \approx 0.980 \quad (15.4)$$

Laboratory demonstrations have achieved 174 K below room temperature in Yb:YLF crystals, making optical refrigeration the only technology to reach cryogenic temperatures without mechanical or thermoelectric systems.

#### Radiation Mode (Thermal Photonics)

Optical refrigeration is inherently a *radiation-mode* process: it converts absorbed photons into higher-energy emitted photons, extracting thermal energy through the radiation field. In the MPM, the optically refrigerated element introduces a radiation-dominated edge ( $\Gamma \gg 10$ ) even in environments where all other paths are conduction-dominated ( $\Gamma \ll 0.1$ ). This mode mismatch is by design—the radiation channel provides the active cooling mechanism.

### Integration Killers: Pumped Cooling

- **TEC:** Parasitic thermal conduction through the module body; contact resistance at solder joints ( $R_c \sim 0.5\text{--}2 \text{ K} \cdot \text{cm}^2/\text{W}$ ); sink saturation under  $Q_{\text{reject}} = P_{\text{cool}}(1 + 1/\text{COP})$ .
- **Optical refrigeration:** Impurity-driven parasitic absorption ( $\eta_{\text{ext}} < \eta_{\text{ext,min}} \Rightarrow$  net heating); optical coupling (need high- $Q$  cavity or waveguide feed without thermal bridges); thermal isolation vs. optical access trade-off (low- $\kappa$  spacer blocks photon delivery).
- **Common:** Active elements draw power  $\Rightarrow$  system-level COP degrades at high  $\Delta T$  or low sink capacity; reliability under thermal cycling (solder fatigue for TEC, optical coating degradation for ASF).

### 15.1.3 MPM Integration: Active Cooling Edges

In the MPM framework, an actively pumped cooler modifies the network in a qualitatively new way: it introduces an edge with an *effective negative resistance* (heat pumped against the gradient). The extended conductance for an active edge is:

$$G_{\text{active}} = G_{\text{passive}} + \frac{P_{\text{cool}}}{\Delta T} \quad (15.5)$$

where  $G_{\text{passive}}$  is the intrinsic (parasitic) conductance and  $P_{\text{cool}}/\Delta T$  is the pumped contribution. When  $P_{\text{cool}} > G_{\text{passive}} \cdot \Delta T$ , the net heat flow reverses direction—the edge acts as a heat pump.

#### Design Rule:

DR-15F.1 — Active Edge Integration When incorporating actively pumped cooling (TEC or optical refrigeration) into an MPM, model the active element as a modified conductance edge:  $G_{\text{active}} = G_{\text{passive}} + P_{\text{cool}}/\Delta T$ . The active element changes the critical path (DR-T.6); recalculate the full network after insertion. Verify that the power budget for the active element ( $P_{\text{elec}}$  for TEC,  $P_{\text{pump}}$  for ASF) does not exceed the system thermal envelope.

### Active Edge Integration Worksheet

Use this 7-step template whenever inserting an active cooling element into an existing thermal network. WE-15.1 below instantiates this worksheet for a TEC in a 3D IC package.

1. **Passive baseline:** Record the existing MPM: all edge conductances  $G_e$ , critical path, and  $\Gamma$  per edge.
2. **Active element specification:** Identify  $P_{\text{cool}}$ ,  $\Delta T_{\text{active}}$ ,  $P_{\text{drive}}$ , and COP (for TEC) or  $\eta_{\text{ext}}$  (for ASF).
3. **Active edge conductance:** Compute  $G_{\text{active}} = G_{\text{passive}} + P_{\text{cool}}/\Delta T$  (Eq. 15.5).
4. **Total heat rejection:**
  - TEC:  $Q_{\text{reject}} = P_{\text{cool}}(1 + 1/\text{COP})$ .
  - ASF:  $Q_{\text{reject}} = P_{\text{cool}} + P_{\text{pump}}(1 - \eta_{\text{ext}})$ .
5. **Sink load update:** Compute new sink temperature rise:  $\Delta T_{\text{sink}} = Q_{\text{reject}}/G_{\text{sink}}$ .
6. **Re-run MPM:** Rebuild the conductance matrix with the active edge inserted. Identify the new critical path. *Does the bottleneck migrate?* (DR-T.6 check.)
7. **Budget check:** Verify  $P_{\text{drive}} + P_{\text{parasitic}} \leq P_{\text{budget}}$ . If FAIL, reduce  $P_{\text{cool}}$  or upgrade sink capacity and return to Step 3.



### 15.1.4 Worked Example: WE-15.1—Active Cooling in a 3D IC Package

#### Case Study:

**WE-15.1: TEC-Augmented 3D IC Stack Problem:** A 3D IC package (logic die + HBM stack on silicon interposer) has a passive MPM with critical path through the TIM to the heat sink ( $G_{\text{TIM}} = 5 \text{ W/K}$ ,  $G_{\text{sink}} = 20 \text{ W/K}$ ). At current power  $P = 150 \text{ W}$ , the passive junction temperature is  $T_j = 25 + 150 \times 0.25 = 62.5^\circ\text{C}$  (within the  $T_j \leq 105^\circ\text{C}$  specification). However, a next-generation die upgrade to  $P = 400 \text{ W}$  on the same package would yield  $T_j = 25 + 400 \times 0.25 = 125^\circ\text{C}$ —exceeding spec by 20 K. A single-stage TEC (COP = 0.5 at  $\Delta T = 30 \text{ K}$ ) is proposed between the TIM and heat sink to recover thermal margin. Evaluate the impact.

#### Solution:

**Step 1: Passive network (Worksheet Step 1).** The critical-path resistance is:

$$R_{\text{passive}} = \frac{1}{G_{\text{TIM}}} + \frac{1}{G_{\text{sink}}} = \frac{1}{5} + \frac{1}{20} = 0.25 \text{ K/W} \quad (15.6)$$

At  $P = 400 \text{ W}$ :  $T_j = 25 + 400 \times 0.25 = 125^\circ\text{C}$ . This exceeds the  $105^\circ\text{C}$  limit by 20 K.

**Step 2: TEC specification (Worksheet Step 2).** Target: reduce  $T_j$  by  $\geq 20 \text{ K}$ . TEC parameters:  $\Delta T_{\text{TEC}} = 30 \text{ K}$ , COP = 0.5.

**Step 3: TEC insertion (Worksheet Steps 3–4).** At COP = 0.5:

$$P_{\text{elec}} = \frac{P_{\text{cool}}}{\text{COP}} = \frac{400}{0.5} = 800 \text{ W} \quad (15.7)$$

Total heat rejected to the sink:

$$Q_{\text{reject}} = P_{\text{cool}} + P_{\text{elec}} = P_{\text{cool}}(1 + 1/\text{COP}) = 400 \times 3 = 1200 \text{ W} \quad (15.8)$$

**Step 4: MPM update (Worksheet Steps 5–6).** The TEC edge conductance:

$$G_{\text{active}} = \frac{P_{\text{cool}}}{\Delta T_{\text{TEC}}} = \frac{400}{30} = 13.3 \text{ W/K} \quad (15.9)$$

However, the sink must now reject 1200 W instead of 400 W. Sink temperature rise:  $\Delta T_{\text{sink}} = 1200/20 = 60 \text{ K}$  above ambient, giving  $T_{\text{sink}} = 85^\circ\text{C}$ . The TEC cold side is 30 K below the hot side:  $T_{\text{cold}} = 85 - 30 = 55^\circ\text{C}$ . With the TIM drop ( $400/5 = 80 \text{ K}$ ):  $T_j = 55 + 80 = 135^\circ\text{C}$ . This is *worse* than the passive case because the  $3\times$  heat multiplication overwhelms the 30 K TEC lift.

**Step 5: Budget check (Worksheet Step 7).**  $P_{\text{elec}} = 800 \text{ W}$  is clearly impractical.  $Q_{\text{reject}} = 1200 \text{ W}$  exceeds typical package-level sink capacity.

**Conclusion:** At COP = 0.5 and  $P = 400 \text{ W}$ , TEC insertion is counterproductive—the sink overload negates the active cooling benefit. The TEC is viable only when (a) the cooling load is small relative to sink capacity, or (b) the COP is significantly higher (multi-stage or advanced thermoelectric materials with  $ZT > 2$ ). DR-15F.1 requires that the total power budget remain within system limits; this example demonstrates why that check is essential.

## 15.2 Phase-Change Reconfigurable Metamaterials

The switchable devices of Chapter 6 (the doublet metadvice, LMTI-based switches, and phase-change thermal diodes) operate at fixed switching thresholds determined by material properties. This section extends those concepts toward *reconfigurable* metamaterials that can continuously

adapt their thermal properties, enabling dynamic MPM topologies.

### 15.2.1 VO<sub>2</sub>: Volatile Switching for Real-Time Control

Vanadium dioxide undergoes a metal-insulator transition at  $T_{MIT} \approx 68^\circ\text{C}$  (Chapter 6, §06.3.3). The key properties for reconfigurable metamaterials are:

- **Thermal conductivity contrast:**  $\kappa_{\text{ins}}/\kappa_{\text{met}} \approx 3.5/6.0 = 0.58$  ( $1.7\times$  switching ratio).
- **Emissivity contrast:** In the LWIR (8–14  $\mu\text{m}$ ), VO<sub>2</sub> emissivity switches from  $\varepsilon_{\text{ins}} \approx 0.9$  to  $\varepsilon_{\text{met}} \approx 0.4$ —a  $2.25\times$  radiative contrast.
- **Switching speed:** The structural phase transition occurs on picosecond timescales when optically triggered, or millisecond timescales for thermally driven transitions.
- **Volatility:** The transition is *volatile*—VO<sub>2</sub> reverts to the insulating phase when the temperature drops below  $T_{MIT}$ . This provides automatic, passive feedback: the material “self-regulates” near the transition temperature.

### 15.2.2 GST: Non-Volatile Switching for Programmable Metamaterials

Ge<sub>2</sub>Sb<sub>2</sub>Te<sub>5</sub> (GST), familiar from optical data storage (DVD-RW, Blu-ray), offers a qualitatively different switching paradigm. Its amorphous-to-crystalline phase transition is *non-volatile*: both phases are stable at room temperature, and switching between them requires an external energy pulse.

**Figure 15.3: Phase-Change Material Properties for Reconfigurable Thermal Metamaterials**




Figure 15.3: Phase-change material properties for reconfigurable thermal metamaterials. (a) Thermal conductivity  $\kappa$  vs. temperature for VO<sub>2</sub> (volatile MIT at  $68^\circ\text{C}$ ,  $1.7\times$  contrast) and GST (non-volatile amorphous-to-crystalline transition,  $\sim 8\times$  contrast). Hysteresis loops are shown for both materials. (b) Emissivity contrast comparison: bulk VO<sub>2</sub> ( $2.25\times$ ), bulk GST ( $2.3\times$ ), and GST metasurface-enhanced ( $4.25\times$ ). The metasurface resonance amplifies the intrinsic emissivity switching by bringing sub-wavelength resonators on/off resonance with the phase change (see Eq. 15.11).

Key properties for thermal metamaterial applications:

Table 15.5: GST phase-change properties for thermal design.

| Property                                   | Amorphous                                 | Crystalline                                |
|--------------------------------------------|-------------------------------------------|--------------------------------------------|
| $\kappa$ (W/m·K)                           | 0.2                                       | 1.5–2.0                                    |
| $n$ at $\lambda = 10 \mu\text{m}$          | 3.5                                       | 6.0                                        |
| $\varepsilon$ (LWIR)                       | 0.3                                       | 0.7                                        |
| Crystallization $T$ ( $^{\circ}\text{C}$ ) | —                                         | 150–170                                    |
| Amorphization pulse                        | —                                         | $>600^{\circ}\text{C}$ , $<100 \text{ ns}$ |
| Cycle endurance                            | $> 10^9$ (optical), $> 10^6$ (electrical) |                                            |

The  $\sim 8\times$  thermal conductivity contrast and  $\sim 2.3\times$  emissivity contrast make GST a far more powerful switching element than  $\text{VO}_2$  for thermal applications. The non-volatility means that a GST-based metamaterial can be *programmed* once and then operates passively, consuming zero standby power.

The emissivity switching of  $\text{VO}_2$  has direct MPM consequences. For a surface at  $T = 350 \text{ K}$  with area  $A = 1 \text{ cm}^2$ :

$$\frac{G_{\text{rad}}^{\text{met}}}{G_{\text{rad}}^{\text{ins}}} = \frac{\varepsilon_{\text{met}}}{\varepsilon_{\text{ins}}} = \frac{0.4}{0.9} \approx 0.44 \quad (15.10)$$

The MPM  $\Gamma$  value for the surface path changes by the same factor, potentially reclassifying the path from mixed-mode to conduction-dominated.

### Multiphysics Integration

GST switching crosses all three mode boundaries simultaneously. In the MPM, a GST element on a path changes: (a)  $G_{\text{cond}}$  by a factor of  $\sim 8$  (thermal conductivity contrast), (b)  $G_{\text{rad}}$  by a factor of  $\sim 2.3$  (emissivity contrast), and (c)  $\Gamma$  by  $\sim 2.3/8 \approx 0.29$  (the radiation fraction decreases in the crystalline state). A path that was mixed-mode ( $\Gamma \approx 0.5$ ) in the amorphous state may become conduction-dominated ( $\Gamma \approx 0.15$ ) after crystallization. This reclassification changes the appropriate design toolkit (DR-T.4).

### 15.2.3 Reconfigurable Emissivity Metasurfaces

Combining phase-change materials with metasurface architectures (Chapter 8) enables emissivity patterns that can be switched pixel-by-pixel, creating a “thermal display” that controls the radiative component of the MPM programmatically.

A representative architecture consists of a GST thin film ( $\sim 100 \text{ nm}$ ) on a metal back-reflector, patterned with sub-wavelength resonators tuned to the atmospheric transparency window ( $8\text{--}13 \mu\text{m}$ ). In the amorphous phase, the resonators are detuned ( $\varepsilon \approx 0.2$ ); after crystallization, the refractive index shift brings them on-resonance ( $\varepsilon \approx 0.85$ ). The emissivity contrast:

$$\frac{\varepsilon_{\text{cryst}}}{\varepsilon_{\text{amorph}}} = \frac{0.85}{0.2} = 4.25 \quad (15.11)$$

This exceeds the intrinsic emissivity contrast of bulk GST ( $2.3\times$ ) because the metasurface resonance amplifies the switching effect. In the  $\Gamma$  calculation, this  $4.25\times$  radiative conductance swing can shift paths between regimes.

**Design Rule:**

DR-15F.2 — Phase-Change Network Reconfiguration When phase-change elements ( $\text{VO}_2$  or GST) are present in a thermal network, the MPM must be evaluated in *both* phase states. If the critical path or its mode classification ( $\Gamma$ ) changes between states, design both configurations to meet thermal requirements independently (DR-T.3 and DR-T.4 applied to each state). Use  $\text{VO}_2$  for self-regulating volatile control near  $T_{\text{MIT}}$ ; use GST for programmed non-volatile reconfiguration.

**Fabrication Challenges**

The principal fabrication challenges for reconfigurable metamaterials in semiconductor applications include: (a) CMOS-compatible deposition of  $\text{VO}_2$  and GST thin films with uniform stoichiometry across 300 mm wafers; (b) integration of electrical or optical switching actuators without introducing parasitic thermal paths; (c) achieving  $> 10^6$  switching cycles for GST without degradation of thermal properties; and (d) maintaining sub-micron feature registration for metasurface patterns under repeated thermal cycling.

## 15.3 Machine Learning in Thermal Design

The MPM workflow (Chapter 4, §04.6) is a systematic six-step procedure: define requirements, construct network, find critical path, classify modes, apply tools, verify and iterate. Machine learning enhances Steps 2–5 by automating the most computationally intensive portions of this workflow.

### 15.3.1 Inverse Design of Metamaterial Structures

Traditional metamaterial design proceeds *forward*: given a geometry, compute the effective properties ( $\kappa_{\parallel}$ ,  $\kappa_{\perp}$ ,  $\varepsilon(\lambda)$ ). Inverse design reverses this: given target properties required by the MPM, find the geometry.

The inverse design problem can be formulated as an optimization:

$$\min_{\mathbf{p}} \mathcal{L}(\mathbf{p}) = \sum_i w_i \left| f_i(\mathbf{p}) - f_i^{\text{target}} \right|^2 + \lambda_{\text{reg}} R(\mathbf{p}) \quad (15.12)$$

where  $\mathbf{p}$  is the design parameter vector (layer thicknesses, fill fractions, material choices, via diameters),  $f_i(\mathbf{p})$  are computed performance metrics ( $\kappa_{\parallel}$ ,  $\kappa_{\perp}$ ,  $\Gamma$ ),  $f_i^{\text{target}}$  are targets from the MPM analysis,  $w_i$  are importance weights, and  $R(\mathbf{p})$  is a regularization term that penalizes fabrication complexity.

Figure 15.4: Forward vs. Inverse Design Workflows for Thermal Metamaterials




Figure 15.4: Forward vs. inverse design workflows for thermal metamaterials. **(a)** Traditional forward design: the engineer specifies geometry, runs FEA simulation, and checks whether the resulting properties ( $\kappa_{\parallel}$ ,  $\kappa_{\perp}$ ,  $\Gamma$ ) meet the MPM targets; iteration is manual and slow. **(b)** ML-assisted inverse design: a neural network surrogate is trained on FEA data, then inverted via gradient descent to find geometries that directly satisfy MPM targets; a fabrication filter enforces manufacturing design rules (DR-5.x, DR-9.x, DR-18.x). The surrogate evaluates in milliseconds, enabling exploration of  $> 10^5$  candidate designs. This parallels the evolution from manual to computational lens design described in *The Eikonal Bridge* (Book 1).

Three ML approaches have shown promise for thermal metamaterial inverse design:

1. **Neural network surrogates.** A forward model  $\mathbf{p} \rightarrow f_i(\mathbf{p})$  is trained on finite-element simulation data, then inverted via gradient descent. Training requires  $10^3$ – $10^4$  forward simulations, after which the surrogate evaluates in milliseconds.
2. **Generative models (GANs, VAEs).** A generative network maps from the target property space directly to a distribution of valid geometries, capturing the one-to-many nature of inverse design (multiple geometries can produce the same effective properties). This is particularly valuable when fabrication constraints eliminate some solutions.
3. **Topology optimization with adjoint methods.** The material distribution is parameterized on a voxel grid and optimized using adjoint-derived gradients from the thermal PDE. This is the most physics-informed approach and naturally respects conservation laws, but is computationally expensive for 3D problems.

### 15.3.2 ML-Optimized Thermal Network Topology

Beyond optimizing individual metamaterial elements, ML can optimize the network topology itself. Given a set of thermal nodes (heat sources, sinks, and interfaces), the network topology—which nodes are connected, through what mode—represents the highest-level design decision in the MPM workflow (Step 2).

**Trilogy Connection:**

ML Inverse Design Across the Trilogy The ML inverse design paradigm in thermal photonics parallels developments in computational lens design (Book 1) and detector optimization (Book 2). In all three domains, the pattern is the same: (a) physics-based forward model, (b) surrogate or adjoint for fast gradients, (c) fabrication constraints as hard boundaries, (d) human-in-the-loop validation. The MPM provides the thermal-domain analog of the merit function in optical design.

**15.3.3 Design Validation: The ML–MPM Interface****Design Rule:**

DR-15F.3 — ML Design Validation Machine learning-generated thermal designs must be validated against the MPM before fabrication. Specifically: (a) verify that all path conductances match targets within 10%, (b) confirm that the critical path identified by ML matches the MPM analysis, (c) check that all fabrication design rules (DR-5.x, DR-9.x, DR-18.x) are satisfied, and (d) run Monte Carlo sensitivity analysis (DR-T.7) on the ML-optimized parameters.

**Integration Killers: ML-Designed Structures**

- **Surrogate fidelity:** Neural network predictions degrade for geometries outside the training distribution; extrapolation errors are silent (no uncertainty estimate without ensemble or Bayesian methods).
- **Fabrication-awareness gap:** ML optimizers exploit geometric freedom that may violate undocumented fab constraints (e.g., aspect-ratio limits, etch lag, alignment tolerances) unless the fabrication design rules (DR-5.x, DR-9.x, DR-18.x) are hard-coded as constraints.
- **Validation cost:** Each ML-proposed design still requires at least one full-physics simulation for DR-15F.3 compliance; the “speed” benefit of ML is in exploration, not in elimination of verification.

### 15.3.4 Worked Example: WE-15.2—ML-Assisted Via Array Optimization

#### Case Study:

**WE-15.2: Neural Network Via Array Design Problem:** A glass interposer (Chapter 7) requires a graded via array to match a non-uniform heat source map with three hotspots. The MPM analysis (Step 4) specifies:

- Path P1 (hotspot 1 to sink):  $G_{\text{cond}}^{\text{target}} = 50 \text{ W/K}$
- Path P2 (hotspot 2 to sink):  $G_{\text{cond}}^{\text{target}} = 30 \text{ W/K}$
- Path P3 (hotspot 3 to sink):  $G_{\text{cond}}^{\text{target}} = 20 \text{ W/K}$
- Total via count budget:  $N_{\text{max}} = 5000$
- Via diameter:  $d_{\text{via}} = 50 \text{ }\mu\text{m}$  (DR-18.1 compliant)
- Substrate:  $100 \text{ }\mu\text{m}$  borosilicate glass ( $\kappa_{\text{glass}} = 1.1 \text{ W/m}\cdot\text{K}$ )

#### Solution:

**Step 1: Surrogate training.** A three-layer fully connected neural network is trained on 5000 COMSOL simulations with randomized via distributions. Input: via density map ( $50 \times 50$  grid). Output:  $G_{\text{cond}}$  for each of the three paths. Training time: 4 hours. Inference time: 2 ms per evaluation.

**Step 2: Inverse optimization.** Gradient descent on the surrogate with fabrication constraints (minimum pitch  $100 \text{ }\mu\text{m}$ , maximum pitch  $500 \text{ }\mu\text{m}$ , total count  $\leq 5000$ ). Convergence in  $\sim 200$  iterations ( $\sim 0.4 \text{ s}$  total).

**Step 3: Result.** The ML-optimized layout clusters vias near hotspots (consistent with DR-9.4, DR-9.5) but also places a low-density perimeter ring that improves lateral heat spreading—a non-obvious design feature that pure rule-based design (Chapter 7) would not produce.

**Step 4: Validation.** Full COMSOL simulation of the ML-proposed layout confirms:

$$G_{\text{P1}} = 52 \text{ W/K}, \quad G_{\text{P2}} = 31 \text{ W/K}, \quad G_{\text{P3}} = 21 \text{ W/K} \quad (15.13)$$

All within 5% of targets. DR-T.6 check: critical path unchanged (P1 remains dominant). DR-T.7 check: Monte Carlo with  $\pm 5\%$  via diameter variation gives  $\sigma_G/G < 3\%$ —within  $3\sigma$  margin.

**Conclusion:** ML inverse design reduces the via array design cycle from days (manual iteration with Chapter 7 rules) to minutes, while respecting all fabrication constraints and producing non-obvious performance improvements. The 5% accuracy confirms DR-15F.3 compliance.

## 15.4 Adaptive Thermal Networks

The ultimate extension of the MPM concept is a network that monitors itself and adjusts its own conductances in real time. This section describes the architecture, control theory, and technology requirements for such systems.

### 15.4.1 Sensor-Driven Conductance Adjustment

An adaptive thermal network consists of three layers (Figure 15.5):

1. **Sensing layer:** Distributed temperature sensors (thermocouples, RTDs, IR imaging,

or NV-center probes from Chapter 14) continuously monitor node temperatures  $\mathbf{T}(t) = [T_1(t), \dots, T_N(t)]^T$ .

2. **Control layer:** A real-time algorithm (PID, model-predictive control, or deep reinforcement learning) computes target conductance adjustments  $\Delta G_e(t)$  for each switchable edge based on the current temperature error  $\mathbf{e}(t) = \mathbf{T}_{\text{target}} - \mathbf{T}(t)$ .
3. **Actuation layer:** Switchable elements (VO<sub>2</sub> surfaces, GST patches, TECs, MEMS shutters, or electronically controlled fans) adjust the physical conductances.

**Figure 15.5: Three-Layer Adaptive Thermal Network Architecture**




Figure 15.5: Three-layer architecture for adaptive thermal networks. The sensing layer (thermocouples, RTDs, NV probes) monitors node temperatures; the control layer (PID, MPC, or DRL) computes conductance adjustments; the actuation layer (VO<sub>2</sub>, GST, TEC, MEMS) modifies physical conductances. The MPM is re-evaluated at each control step to track critical-path migration (DR-T.6).

The mathematical framework for adaptive networks builds on the transient thermal network equation from Chapter 2:

$$\mathbf{C} \frac{d\mathbf{T}}{dt} + \mathbf{G}(t)\mathbf{T} = \mathbf{Q}(t) \quad (15.14)$$

where  $\mathbf{G}(t)$  is now *time-varying* because the switchable elements change their conductances under control.

### 15.4.2 Control Theory for Thermal Networks

The control system must operate slower than the fastest thermal mode to avoid instability. The Nyquist criterion adapted for thermal control requires:

$$\tau_{\text{control}} > 2 \tau_{\text{th,min}} \quad (15.15)$$

For typical semiconductor packaging elements:



Table 15.6: Thermal time constants and control bandwidth limits for representative elements.

| Element                                      | $C$ (J/K)  | $G$ (W/K) | $\tau_{th}$       |
|----------------------------------------------|------------|-----------|-------------------|
| Thin-film VO <sub>2</sub> (1 $\mu\text{m}$ ) | $10^{-8}$  | 0.01      | 1 $\mu\text{s}$   |
| Silicon die (0.5 mm)                         | 0.8        | 50        | 16 ms             |
| Copper heat sink                             | 38         | 20        | 1.9 s             |
| GST pixel (10 $\times$ 10 $\mu\text{m}$ )    | $10^{-11}$ | $10^{-4}$ | 0.1 $\mu\text{s}$ |
| TEC module                                   | 15         | 5         | 3 s               |

### Steady-State Limitation

The steady-state limitation (Chapter 2, DR-02.1) applies to adaptive networks: the control system cannot violate the maximum principle. It can redistribute heat flux and shift the MPM critical path, but it cannot reduce the temperature below  $T_{\text{sink}}$  without an active cooler (§15.1). The combination of adaptive network + active cooling is the most powerful but also the most complex configuration.

### Integration Killers: Adaptive Thermal Networks

- **Sensor placement and density:** Temperature sensor count and placement must resolve the thermal modes of interest; under-sampling aliases high-frequency transients into erroneous control signals.
- **Control latency:** Sensor-to-actuator loop delay must satisfy  $\tau_{\text{loop}} < \tau_{th,\text{min}}/2$  (DR-15F.4); for thin-film VO<sub>2</sub> elements ( $\tau_{th} \sim 1 \mu\text{s}$ ), this requires sub-microsecond control—challenging for embedded processors.
- **Fail-safe reliability:** All actuators must revert to a thermally safe passive state on control loss; latching (non-volatile) actuators like GST hold their last state, which may not be the safe state—requires explicit reset circuitry.
- **System qualification:** No existing semiconductor reliability standard (JEDEC, AEC-Q) covers adaptive thermal networks; qualification must be developed per DR-15F.4 criteria.

### 15.4.3 Autonomous Thermal Management

Deep reinforcement learning (DRL) offers a path to autonomous thermal management systems that learn optimal control policies through interaction with the thermal environment. The DRL agent observes the temperature state, takes actions (adjusting switchable conductances), and receives rewards based on temperature uniformity and power efficiency.

The key advantage of DRL over model-predictive control is that it can learn non-obvious strategies from experience. For example, a DRL agent managing a 3D IC package with five switchable elements discovered that briefly *increasing* the conductance of a non-critical path during a power transient improved the critical-path temperature by pre-distributing heat—a strategy that violates the steady-state intuition of “focus on the critical path” (DR-T.3) but exploits transient dynamics (Chapter 2).

**Design Rule:**

DR-15F.4 — Adaptive Network Stability Adaptive thermal networks must satisfy: (a) control bandwidth  $< 1/(2\tau_{\text{th,min}})$  to avoid thermal instability, (b) fail-safe default state (all actuators revert to a thermally safe passive configuration on control loss), and (c) MPM re-evaluation at every control step to track critical path shifts (DR-T.6). The passive fallback design must independently meet thermal requirements with appropriate margin (DR-T.7).

## 15.5 Thermal Computing: A Speculative Horizon

*This section is a pedagogical horizon scan, not a design recommendation. We include it to map the boundary between what the MPM can address today and what remains beyond current reach, and to provide the MPM-based feasibility metrics (Table 15.7) that justify this boundary.*

The preceding sections describe technologies with clear paths to production relevance within 5–10 years. This section examines a far more speculative frontier: using heat flow itself as a computational medium. We include this topic for completeness and intellectual context, with a clear feasibility assessment.

### 15.5.1 Heat-Based Logic Gates

Thermal diodes (Chapter 6) provide the nonlinearity required for logic. A thermal AND gate can be constructed from two diodes in series: both temperature inputs must exceed the threshold for the output to be “high.” Similarly, thermal transistors (three-terminal devices where a gate temperature controls the source-drain conductance) have been demonstrated using negative differential thermal resistance in phononic systems.

The fundamental metrics for thermal logic are summarized in Figure 15.6:


**Figure 15.6: Thermal Logic vs. Electronic (CMOS) Logic - Fundamental Comparison**

Figure 15.6: Log-scale comparison of fundamental performance metrics for thermal logic gates vs. electronic (CMOS) logic. Four metrics are compared: signal propagation speed (m/s), switching energy per gate (J), gate areal density (/cm<sup>2</sup>), and effective fan-out. Thermal logic is  $10^4$ – $10^{14}$  times inferior across all metrics. The signal attenuation bar (rightmost) highlights the critical cascadability limitation: thermal signals decay as  $1/\sqrt{t}$  in diffusive media with no regeneration mechanism, whereas electronic transistors provide full signal restoration.

Table 15.7: Thermal vs. electronic logic: fundamental comparison.

| Metric           | Thermal Logic                    | Electronic (CMOS)                   |
|------------------|----------------------------------|-------------------------------------|
| Signal speed     | $10^{-6}$ – $10^0$ m/s           | $\sim 10^8$ m/s                     |
| Switching energy | $10^{-12}$ – $10^{-6}$ J         | $10^{-18}$ – $10^{-15}$ J           |
| Gate density     | $10^2$ – $10^4$ /cm <sup>2</sup> | $10^8$ – $10^{10}$ /cm <sup>2</sup> |
| Fan-out          | 2–3                              | >10                                 |
| Cascadability    | Poor (signal attenuation)        | Excellent (regenerative)            |

The comparison is sobering. Thermal signals propagate  $10^8$ – $10^{14}$  times slower than electrical signals, consume  $10^3$ – $10^9$  times more energy per switching event, and achieve  $10^4$ – $10^8$  times lower gate density. The signal attenuation problem (thermal signals decay as  $1/\sqrt{t}$  in diffusive media) is particularly severe: there is no thermal equivalent of the electronic transistor’s signal regeneration.

### 15.5.2 Thermal Memory

GST-based thermal memory is more promising than thermal logic because memory requires only *bistability*, not *cascadability*. A GST element programmed to its amorphous or crystalline state stores one bit of thermal information that manifests as a conductivity contrast. Reading the state requires measuring  $\kappa$  or  $\varepsilon$ —achievable with the metrology tools of Chapter 13.

### 15.5.3 Feasibility Assessment

Thermal computing as a *replacement* for electronic computing is not viable: the fundamental physics gaps are  $> 10^4 \times$  across every performance metric. However, thermal computing as a *co-processor* for specific niche applications—particularly in extreme environments where electronics

cannot operate (inside nuclear reactors, near volcanic vents, in certain space environments)—remains a legitimate research topic.

## 15.6 Closing the Trilogy

### 15.6.1 The Photonic Engineering Design Stack

The three books of the Open-Source Engineering Optics Trilogy form a complete photonic engineering design stack:

**Figure 15.7: The Photon Energy Lifecycle Across the Trilogy**




Figure 15.7: The Open-Source Engineering Optics Trilogy design stack. Book 1 (*The Eikonal Bridge*) covers propagation; Book 2 (*Quantum Field Imaging*) covers detection; Book 3 (*Thermal Photonics*) covers energy management. Together they span the complete lifecycle of a photon in any engineered system.

#### Trilogy Connection:

##### The Complete Photon Lifecycle

1. **Propagation** (Book 1, *The Eikonal Bridge*): How photons travel through optical systems—ray tracing, wave optics, phase-space methods. The duality: rays vs. waves.
2. **Detection** (Book 2, *Quantum Field Imaging*): How photons are converted to information—coherence, quantum noise, detector architectures. The duality: classical vs. quantum fields.
3. **Management** (Book 3, *Thermal Photonics*): How thermal energy (phonons and photons) is controlled—conduction engineering, radiation design, network optimization. The duality: phonon diffusion vs. photon emission.

### 15.6.2 Grand Challenges

Five fundamental challenges remain at the intersection of thermal photonics and semiconductor technology:

1. **The 3D integration thermal barrier:** As die stacking exceeds 8 layers, the interior dies have no direct access to heat sinks. MPM analysis shows that interior critical paths

are radiation-dominated at  $T > 400$  K—a regime poorly served by current conduction tools.

2. **The athermal optic:** Combining the Fourier–Snell bridge (Chapter 1) with metamaterial thermal control to create truly athermal optical elements whose performance is invariant to thermal environment.
3. **The quantum–thermal interface:** Bridging Chapter 14’s quantum thermal signatures with production thermal management to enable quantum-limited thermal diagnostics in semiconductor fabs.
4. **The autonomous thermal fab:** Extending §15.4 to facility-scale adaptive thermal networks where the entire fab operates as a single optimized thermal system.
5. **The thermal Rosetta Stone:** Developing a universal translation layer between thermal, electrical, optical, and mechanical design domains, with the MPM as the common language.

### 15.6.3 The Thermal Photonics Manifesto

#### The Thermal Photonics Manifesto

Heat is not waste—it is information encoded in a temperature field, waiting to be read, redirected, and engineered. The Mode-Path Matrix is the Rosetta Stone that translates between the language of heat and the language of design. Every thermal problem is a network problem; every network problem has a critical path; every critical path has a dominant mode. Find the mode, fix the path, verify the network. This is the discipline of thermal photonics.

These challenges span the space from near-term engineering (thermal-optical co-design) to long-term fundamental research (quantum-thermal interface). The MPM provides the common language for all of them: regardless of the specific physics, every thermal system is a network of conductances with modes and critical paths. The framework scales from millikelvin cryostats (Chapter 14) to high-temperature semiconductor processes (Chapter 11), and from passive metamaterials (Chapters 5–9) to adaptive intelligent networks (this chapter).

## Technology Readiness Summary

Table 15.8 and Figure 15.8 summarize the technology readiness of all active and adaptive methods discussed in this chapter, evaluated against semiconductor production requirements.


**Figure 15.8: Technology Readiness Landscape for Active Thermal Control**

Figure 15.8: Technology readiness landscape for active thermal control. Each bubble represents a technology from Table 15.8, plotted by current TRL (vertical axis) vs. estimated timeline to semiconductor production readiness (horizontal axis). Bubble color indicates the  $\Gamma$  impact category: blue (conduction), red (radiation), purple (both modes), green (design tool, mode-independent). Bubble size scales with estimated production impact. The shaded green region marks the “production-ready zone” ( $\text{TRL} \geq 7$ ,  $\text{timeline} \leq 3$  years). Only TEC modules and ML inverse design currently occupy this zone.

Table 15.8: Technology readiness assessment for active thermal control methods. The  $\Delta\Gamma$  column (v4 addition) quantifies the typical magnitude of the mode-dominance ratio swing and whether the technology can cause a regime crossing ( $\Gamma$  crossing the 0.1 or 10 boundary).

| Technology                | TRL | $\Gamma$ | Typical<br>Regime Crossing                                                                | $\Delta\Gamma$ /                        | Timeline | MPM Role                       |
|---------------------------|-----|----------|-------------------------------------------------------------------------------------------|-----------------------------------------|----------|--------------------------------|
| TEC modules               | 9   | Cond.    | $G_{\text{eff}}$ can become negative; stays conduction-dominated                          | be-<br>negative;<br>conduction-         | Now      | Active negative- $R$ edge      |
| Optical refrig.           | 4   | Rad.     | Injects $\Gamma \gg 10$ into cond. network                                                | edge                                    | 5–10 yr  | Radiation-mode active edge     |
| VO <sub>2</sub> switching | 5   | Both     | $\varepsilon$ : 2.3 $\times$ ; can cross $\Gamma = 1$ boundary                            | can cross                               | 3–5 yr   | Volatile dynamic edge          |
| GST reconfig.             | 4   | Both     | $\kappa$ : 8 $\times$ , $\varepsilon$ : 4.25 $\times$ (metasurf.); regime crossing likely | 4.25 $\times$<br>regime crossing likely | 5–8 yr   | Non-volatile programmable edge |
| ML inverse design         | 6   | N/A      | Targets any $\Gamma$ ; no intrinsic swing                                                 |                                         | 1–3 yr   | Automated Step 3/4 of workflow |
| ML topology opt.          | 4   | N/A      | Can restructure topology; $\Delta\Gamma$ problem-dependent                                | $\Delta\Gamma$                          | 3–5 yr   | Automated Step 2 of workflow   |
| Adaptive networks         | 3   | Both     | Continuous $\Gamma$ tuning via switchable edges                                           |                                         | 5–10 yr  | Closed-loop dynamic MPM        |
| Thermal logic             | 2   | Cond.    | $\Delta\kappa \sim 3\times$ (thermal diode); stays conduction-dominated                   | (thermal diode);<br>stays conduction-   | >15 yr   | Speculative                    |
| Thermal memory            | 4   | Both     | Same as GST reconfiguration                                                               |                                         | 5–8 yr   | Programmable network state     |

Table 15.9: Complete set of Chapter 15 design rules (DR-15F series).

| Rule     | Section | Description                                                                                                                                           |
|----------|---------|-------------------------------------------------------------------------------------------------------------------------------------------------------|
| DR-15F.1 | §15.1.3 | Active edge integration: model pumped cooling as modified conductance; recalculate MPM after insertion; verify total power budget.                    |
| DR-15F.2 | §15.2.3 | Phase-change reconfiguration: evaluate MPM in both phase states; design for thermal requirements in each state independently.                         |
| DR-15F.3 | §15.3.3 | ML design validation: verify ML-generated designs against MPM targets, fabrication DRs, and Monte Carlo sensitivity before committing to fabrication. |
| DR-15F.4 | §15.4.3 | Adaptive network stability: control bandwidth $< 1/(2\tau_{\text{th,min}})$ ; passive fail-safe required; MPM re-evaluation at every control step.    |

## Looking Forward

This is the final chapter of the main text. The tools developed across fifteen chapters—from the fundamental duality of phonons and photons (Chapter 1) through the MPM framework (Chapter 4), the conduction toolkit (Chapters 5–7), the radiation toolkit (Chapters 8–10),

semiconductor applications (Chapters 11–13), the quantum-thermal frontier (Chapter 14), and adaptive thermal intelligence (this chapter)—constitute a complete, systematic methodology for thermal engineering in semiconductor fabrication and beyond.

The appendices that follow provide reference material: Appendix A collects the mathematical tools, Appendix B tabulates material properties, Appendix C provides design worksheets for routine calculations, and Appendix D treats convection as a boundary condition. Together with the companion Python simulation code available on the project GitHub repository, these resources aim to make the MPM methodology immediately actionable for practicing engineers.

The trilogy is complete. The photonic engineering design stack—propagation, detection, management—provides the intellectual framework. The Mode-Path Matrix provides the practical methodology. The rest is engineering.



## Problems

- P-15.1 (Open-ended research)** A 3D IC package with four die layers generates a non-uniform heat map that shifts between two operating modes (compute-intensive and memory-intensive). Design a reconfigurable thermal network using GST-based switchable elements that optimizes the MPM for both modes independently. Your design should specify: (a) the number and placement of switchable elements, (b) the GST programming protocol, (c) the MPM for each mode (including  $\Gamma$  and critical path), and (d) the fail-safe passive configuration. Assess feasibility using current fabrication capabilities (Chapter 13).
- P-15.2 (Open-ended research)** The “thermal Rosetta Stone” concept (Section 15.6.3) proposes five grand challenges. Select one and develop a 3-page research proposal that includes: (a) a specific hypothesis testable within 3 years, (b) the MPM formulation of the problem, (c) the required experimental or computational tools, (d) expected deliverables with quantitative success metrics, and (e) connections to at least two other chapters in this book. Identify which design rules from the DR-T series apply and whether new design rules would be needed.
- P-15.3 (Open-ended research)** Consider a semiconductor fabrication facility (“fab”) as a single thermal system with  $\sim 100$  process tools, each generating 10–100 kW of heat, connected by facility cooling water, cleanroom air handling, and exhaust systems. Construct a facility-level MPM (at the tool-to-tool scale, not the chip scale) and identify: (a) the critical thermal path in the facility, (b) whether the facility  $\Gamma$  is conduction-dominated (water cooling) or radiation-dominated (waste heat to environment), (c) opportunities for waste heat recovery using TPV concepts from Chapter 10, and (d) how adaptive network concepts from §15.4 could be applied at the facility scale. This problem has no unique answer; the goal is to demonstrate that the MPM framework scales from nanometers to buildings.
- P-15.4 (Active edge analysis)** Using the Active Edge Integration Worksheet (§15.1.3), evaluate the insertion of an optical refrigerator (Yb:YLF crystal,  $\eta_{\text{ext}} = 0.996$ , pump power  $P_{\text{pump}} = 10$  W, cooling power  $P_{\text{cool}} = 0.2$  W) into a cryogenic detector package (Chapter 14). The passive MPM has  $G_{\text{cond}} = 0.01$  W/K (mechanical mounts) and  $G_{\text{rad}} = 0.001$  W/K (radiation shields) from the detector (at 80 K) to the cold plate (at 50 K). (a) Compute  $G_{\text{active}}$  and the new critical path. (b) Does  $\Gamma$  change regime? (c) What is the total heat load on the cold plate?
- P-15.5 (Phase-change MPM)** A thermal path in a wafer chuck contains a  $\text{VO}_2$  thin-film element. At  $T = 60^\circ\text{C}$  (below  $T_{\text{MIT}}$ ), the path has  $G_{\text{cond}} = 2.0$  W/K and  $G_{\text{rad}} = 0.5$  W/K. At  $T = 80^\circ\text{C}$  (above  $T_{\text{MIT}}$ ), both conductances change as described in §15.2.1. (a) Compute  $\Gamma$  in both states. (b) Does the mode classification change? (c) If this path is the critical path, what does DR-15F.2 require you to do?

## References by Category

References are organized by theme to serve as a reading guide for further exploration of each frontier topic.

### A. Active Cooling: Thermoelectric and Optical Refrigeration

*This category covers externally pumped thermal control—from mature thermoelectric devices to the frontier of all-solid-state optical cryocoolers based on anti-Stokes fluorescence. Start with [A1] for TEC fundamentals, then [A3–A5] for the physics and state of the art in laser cooling of solids, and [A8] for the landmark all-solid-state cryocooler demonstration.*

- [A1] G. S. Nolas, J. Sharp, and H. J. Goldsmid, *Thermoelectrics: Basic Principles and New Materials Developments*, Springer, 2001.
- [A2] G. J. Snyder and E. S. Toberer, “Complex thermoelectric materials,” *Nat. Mater.*, vol. 7, pp. 105–114, 2008.
- [A3] R. I. Epstein and M. Sheik-Bahae, eds., *Optical Refrigeration: Science and Applications of Laser Cooling of Solids*, Wiley-VCH, 2009.
- [A4] D. V. Seletskiy, R. I. Epstein, and M. Sheik-Bahae, “Laser cooling in solids: advances and prospects,” *Rep. Prog. Phys.*, vol. 79, p. 096401, 2016.
- [A5] M. Sheik-Bahae and R. I. Epstein, “Optical refrigeration,” *Nat. Photon.*, vol. 1, pp. 693–699, 2007.
- [A6] S. D. Melgaard, A. R. Albrecht, M. P. Hehlen, and M. Sheik-Bahae, “Solid-state optical refrigeration to sub-100 Kelvin regime,” *Sci. Rep.*, vol. 6, p. 20380, 2016.
- [A7] R. I. Epstein, M. I. Buchwald, J. Edwards, T. R. Gosnell, and C. E. Mungan, “Observation of laser-induced fluorescent cooling of a solid,” *Nature*, vol. 377, pp. 500–503, 1995.
- [A8] M. P. Hehlen *et al.*, “First demonstration of an all-solid-state optical cryocooler,” *Light Sci. Appl.*, vol. 7, p. 15, 2018.
- [A9] J. B. Khurgin, “Photoluminescent cooling with incoherent light,” *APL Photon.*, vol. 9, p. 081301, 2024.

### B. Phase-Change Reconfigurable Materials (VO<sub>2</sub>, GST)

*Vanadium dioxide and chalcogenide phase-change alloys enable thermally, electrically, or optically triggered switching of both thermal conductivity and emissivity. Begin with [B1–B2] for the MIT physics, [B4–B5] for GST fundamentals, and [B6–B7] for thermal-specific switching demonstrations.*

- [B1] F. J. Morin, “Oxides which show a metal-to-insulator transition at the Néel temperature,” *Phys. Rev. Lett.*, vol. 3, pp. 34–36, 1959.
- [B2] Z. Shao, X. Cao, H. Luo, and P. Jin, “Recent progress in the phase-transition mechanism and modulation of vanadium dioxide materials,” *NPG Asia Mater.*, vol. 10, pp. 581–605, 2018.
- [B3] C. Wen *et al.*, “A review of the preparation, properties and applications of VO<sub>2</sub> thin films with the reversible phase transition,” *Front. Mater.*, vol. 11, p. 1341518, 2024.

- [B4] M. Wuttig and N. Yamada, “Phase-change materials for rewriteable data storage,” *Nat. Mater.*, vol. 6, pp. 824–832, 2007.
- [B5] P. Noé *et al.*, “Phase-change materials for non-volatile memory devices: from technological challenges to materials science issues,” *Semicond. Sci. Technol.*, vol. 33, p. 013002, 2018.
- [B6] K. Ito *et al.*, “Radiative thermal switching via multi-body effects in vanadium-dioxide-based near-field systems,” *Appl. Phys. Lett.*, vol. 105, p. 253503, 2014.
- [B7] C. Fan, H. Xie, and J. Wang, “Realization of bifunctional photothermal and radiative cooling based on self-adaptive VO<sub>2</sub>/PDMS metamaterial,” *Sol. Energy*, vol. 266, p. 112192, 2023.

### C. Machine Learning for Thermal and Photonic Design

*ML-based inverse design accelerates metamaterial development by orders of magnitude. Reference [C1] is the definitive review for thermal metamaterials; [C3–C4] cover photonic inverse design; [C6–C7] address topology optimization and differentiable physics.*

- [C1] H. Wei *et al.*, “Machine learning aided design and optimization of thermal metamaterials,” *Chem. Rev.*, vol. 124, pp. 4881–4933, 2024.
- [C2] J. Jiang, D. Sell, S. Hoyer, J. Hickey, J. Yang, and J. A. Fan, “Free-form diffractive metagrating design based on generative adversarial networks,” *ACS Nano*, vol. 13, pp. 8872–8878, 2019.
- [C3] W. Ma, F. Cheng, and Y. Liu, “Deep-learning-enabled on-demand design of chiral metamaterials,” *ACS Nano*, vol. 12, pp. 6326–6334, 2018.
- [C4] S. So, T. Badloe, J. Noh, J. Bravo-Abad, and J. Rho, “Deep learning enabled inverse design in nanophotonics,” *Nanophotonics*, vol. 9, pp. 1041–1057, 2020.
- [C5] J. Jiang, M. Chen, and J. A. Fan, “Deep neural networks for the evaluation and design of photonic structures,” *Nat. Rev. Mater.*, vol. 6, pp. 679–700, 2021.
- [C6] M. P. Bendsøe and O. Sigmund, *Topology Optimization: Theory, Methods, and Applications*, 2nd ed., Springer, 2004.
- [C7] A. Y. Piggott, J. Lu, K. G. Lagoudakis, J. Petykiewicz, T. M. Babinec, and J. Vučković, “Inverse design and demonstration of a compact and broadband on-chip wavelength demultiplexer,” *Nat. Photon.*, vol. 9, pp. 374–377, 2015.
- [C8] T. W. Hughes, M. Minkov, I. A. D. Williamson, and S. Fan, “Adjoint method and inverse design for nonlinear nanophotonic devices,” *ACS Photon.*, vol. 5, pp. 4781–4787, 2018.

### D. Adaptive Thermal Networks and Control

*Adaptive thermal management extends the MPM from a static design tool to a real-time control framework. [D1–D2] cover model-predictive control for thermal systems, [D3] addresses digital-twin integration, and [D4–D5] review sensor-actuator architectures for active thermal routing.*

- [D1] F. Borrelli, A. Bemporad, and M. Morari, *Predictive Control for Linear and Hybrid Systems*, Cambridge University Press, 2017.
- [D2] Y. Ma, F. Borrelli, B. Hency, B. Coffey, S. Benga, and P. Haves, “Model predictive control for the operation of building cooling systems,” *IEEE Trans. Control Syst. Technol.*, vol. 20, pp. 796–803, 2012.

- [D3] Q. Lu, X. Xie, A. J. Parlikad, and J. M. Schooling, “Digital twin-enabled anomaly detection for built asset monitoring in operation and maintenance,” *Autom. Constr.*, vol. 118, p. 103277, 2020.
- [D4] R. Schacht *et al.*, “Adaptive thermal management for portable electronic devices,” *Appl. Energy*, vol. 346, p. 121387, 2023.
- [D5] M. Schmidt, M. Assadi, and D. Bruch, “Smart thermal management systems: state of the art and future challenges,” *Energy Convers. Manag.*, vol. 267, p. 115899, 2022.

## E. Thermal Computing and Phononics

*Thermal computing explores heat-based logic and memory as speculative analogs to electronic circuits. [E1] is the foundational proposal; [E2–E3] cover phononic crystal engineering; [E4–E5] address the fundamental speed and efficiency limits that constrain thermal processors.*

- [E1] N. Li, J. Ren, L. Wang, G. Zhang, P. Hänggi, and B. Li, “Colloquium: Phononics: manipulating heat flow with electronic analogs and beyond,” *Rev. Mod. Phys.*, vol. 84, pp. 1045–1066, 2012.
- [E2] M. Maldovan, “Sound and heat revolutions in phononics,” *Nature*, vol. 503, pp. 209–217, 2013.
- [E3] M. Maldovan, “Phonon wave interference and thermal bandgap materials,” *Nat. Mater.*, vol. 14, pp. 667–674, 2015.
- [E4] L. Wang and B. Li, “Thermal logic gates: computation with phonons,” *Phys. Rev. Lett.*, vol. 99, p. 177208, 2007.
- [E5] L. Wang and B. Li, “Thermal memory: a storage of phononic information,” *Phys. Rev. Lett.*, vol. 101, p. 267203, 2008.
- [E6] B. Li, L. Wang, and G. Casati, “Thermal diode: rectification of heat flux,” *Phys. Rev. Lett.*, vol. 93, p. 184301, 2004.

## F. Thermal Metamaterials — Foundational Reviews

*These reviews span the full landscape of thermal metamaterials and connect to the book’s core framework. [F1] is the definitive Nature Reviews survey; [F2] covers transformation thermotics foundations; [F3–F5] address near-field, nano/microscale, and phonon-engineering perspectives.*

- [F1] Y. Li *et al.*, “Transforming heat transfer with thermal metamaterials and devices,” *Nat. Rev. Mater.*, vol. 6, pp. 488–507, 2021.
- [F2] S. Xu, H. Xu, and J. Huang, *Transformation Thermotics and Extended Theories: Inside and Outside Metamaterials*, Springer, 2023.
- [F3] Z. Zhang, *Nano/Microscale Heat Transfer*, 2nd ed., Springer, 2020.
- [F4] G. Chen, *Nanoscale Energy Transport and Conversion*, Oxford University Press, 2005.
- [F5] S. Fan, “Thermal photonics and energy applications,” *Joule*, vol. 1, pp. 264–273, 2017.

## G. Trilogy Cross-References

- [G1] J.-L. Chern, *The Eikonal Bridge: A Unified Framework for Ray and Wave Optics*. Open-Source Engineering Optics Trilogy, Book 1, 2025.
- [G2] J.-L. Chern, *Quantum Field Imaging: From Classical Coherence to Quantum Detection*. Open-Source Engineering Optics Trilogy, Book 2, 2025.

---

*Chapter 15 of “Thermal Photonics: A Unified Framework for Engineering Heat Flux”*

*Book 3 of the Open-Source Engineering Optics Trilogy*

*Cross-references: Book 1 (The Eikonal Bridge) — ML inverse design parallels computational lens optimization;*

*Book 2 (Quantum Field Imaging) — NV-center probes for adaptive network sensing (Ch. 14 connection)*